

Recursive Genomic Q-Former: Parameter-Reduction vs. Performance Trade-off

Task details	
Task	cancer_type — 4-way tumor-of-origin classification from bulk gene-expression profiles (TCGA RNA-seq).
Dataset	Unified TCGA cohort: 2,738 samples $\times$ 20,530 genes; top 4,000 highly-variable genes used as input tokens.
Classes	Breast (BRCA, 867), Head-neck (HNSC, 423), Lung (LUNG, 943), Prostate (PRAD, 505).
Split	70/15/15 train/val/test, 523 test samples , seed 42.
Model	Recursive Genomic Q-Former: $d_{\text{model}} = 128$, 256 marker query slots, recursion depth 4.
Metric	Test accuracy and macro-averaged F1.




Figure 1: Pipeline. The shared transformer block (blue) is reused across recursion steps; MoR routes tokens to a per-token depth instead of a fixed pass.

All results are on the **cancer_type** classification head (4 classes: BRCA, HNSC, LUNG, PRAD; 523 test samples), with a fixed recursion depth of 4, $d_{\text{model}} = 128$, 256 marker query slots, and seed 42. The “regular” recursive model shares one transformer block across all recursion steps; the comparison models either un-tie those weights (*Independent*) or add token-adaptive routing (*MoR*).

Table 1: Parameters and performance for all three model types. One view of the full trade-off: parameter budget (recurrent-block params, total params, and the independent/shared ratio) alongside all three performance metrics (accuracy, macro-F1, weighted-F1). Sharing the transformer block cuts recurrent params by $\sim 4\times$ for a 0.57-point accuracy cost; Mixture-of-R recursions (MoR) keeps that small budget yet recovers the gap to 0.19 points and is best on macro-F1.

Model	Weights	Routing	Parameters			Performance		
			Recurrent	Total	Ratio	Acc.	Macro-F1	Wtd-F1
Independent stack	untied	none	530,432	1,109,897	$3.99\times$	0.9904	0.9891	0.9904
Recursive (shared)	shared	none	132,992	712,457	$1.00\times$	0.9847	0.9871	0.9847
MoR (token-adaptive)	shared	token	132,992	712,457	$1.00\times$	0.9885	0.9892	0.9885

Read-out. The shared recursive model uses **74.9% fewer** transformer parameters than the independent stack (132,992 vs. 530,432) for a **0.57-point** accuracy cost. Adding MoR token routing

closes that gap to **0.19 points** at the same small parameter budget, and is the only configuration that also reduces FLOPs.

Table 2: Parameter scaling vs. recursion depth. Shared-weight recursion keeps the transformer parameter count essentially flat as depth grows, whereas an independent stack grows linearly. The reduction factor reaches nearly 8 $\times$ at depth 8.

Recursion depth	Shared params	Independent params	Reduction factor
1	132,608	132,608	1.00 $\times$
2	132,736	265,216	2.00 $\times$
4	132,992	530,432	3.99 $\times$
6	133,248	795,648	5.97 $\times$
8	133,504	1,060,864	7.95 $\times$

Table 3: Compute trade-off (depth-4 variants). Early-exit / adaptive recursion lowers the *effective* mean depth and FLOPs per sample relative to a fixed depth-4 pass, at equal parameter count. MoR achieves the lowest compute and the best F1.

Configuration	Transf. params	Mean depth	MFLOPs/sample	Test acc.
Fixed depth-4 (no early exit)	132,480	4.00	268.4	0.9847
Recursive (shared, early-exit)	132,992	2.75	161.5	0.9847
No recursive refine	132,992	2.75	161.5	0.9771
MoR (token-adaptive)	132,992	2.69	159.1	0.9885

Notes. Param ratio in Table 1 is independent/shared transformer parameters. “Transf. params” counts only the recurrent transformer block (the full model incl. embeddings/heads is $\sim$ 712K for shared variants, $\sim$ 1.11M for the independent stack). Compute-saving for the early-exit and MoR variants is $\approx$ 40% of the fixed depth-4 FLOP budget. Token-reduction at the marker stage additionally compresses 4000 input genes to 256 retained markers (93.6% token reduction) at 98.5% accuracy.

Dataset tasks and per-task results

The model is evaluated on the **unified TCGA benchmark**: 2,738 patient samples, 20,530 genes (top 4,000 highly-variable genes used as input tokens). The same recursive architecture (shared block, depth 4, 256 markers) is trained per task. Six prediction targets span tumor identity, clinical staging, and outcome.

Table 4: Per-task accuracy of the three model types, with parameters. Test accuracy for each of the three models (Independent / Recursive-shared / MoR) across all six TCGA tasks; recurrent-block parameters are listed under each model. The shared recursive model matches the $\sim 4\times$ -larger Independent stack within a couple of accuracy points on most tasks — and even edges it on `pathologic_stage` and `tumor_status` — at one quarter the recurrent parameters.

Task	Cls.	Test n	Test accuracy by model		
			Independent 530,432 par.	Recursive (shared) 132,992 par.	MoR 132,992 par.
<code>cancer_type</code>	4	523	0.9904	0.9847	0.9885
<code>pathologic_stage</code>	4	548	0.4909	0.5000	0.4745
<code>pathologic_T</code>	4	548	0.4653	0.4580	0.4288
<code>pathologic_N</code>	4	548	0.4672	0.4416	0.4124
<code>os_binary</code>	2	548	0.7573	0.7372	0.6350
<code>tumor_status</code>	2	548	0.6734	0.6825	0.6296

Notes. Parameters under each header are recurrent transformer-block parameters (total params: 1,109,897 Independent; 712,457 shared and MoR). All numbers are single-seed (seed 42) test accuracies; the `cancer_type` heads come from `results/{independent,main,mor_token}.json` and the five staging/survival tasks from `results_genonet{independent,,mor}/`. All three model types were trained on every task; the page-1 table reports the same three-way comparison with full parameter and multi-metric detail on `cancer_type`.